




Full 20 MW installed GPT-OSS-20B / A100 fleet; two equal destinations at 50% load  
Assumed shared WAN budget 100 Gbit/s; node endpoint ceilings also apply

Removed source workload (%)



Experimental: regional timing validation fails; long-context batches use serialization fallback. Bands span four snapshots  $\times$  nine calibration/network cases.

