



Full 20 MW installed GPT-OSS-20B / A100 fleet; two equal destinations at 50% load
Assumed shared WAN budget 40000 Gbit/s; node endpoint ceilings also apply

Removed source workload (%)


Original coding contexts; 8 sessions/GPU



24K-32K contexts; 8 sessions/GPU



24K-32K contexts; 16 sessions/GPU



Experimental: regional timing validation fails; long-context batches use serialization fallback. Bands span four snapshots $\times$ nine calibration/network cases.

Queue-Haul LP Queue-Haul Greedy KV Migrate Only Replay Context Only Isolated Fastest