

GPT-OSS / A100: 64-GPU source pool, two 64-GPU destinations, 50% resident load
8 GPUs/node; 512 source sessions; assumed 40 Gbit/s shared WAN

Queue-Haul LP Queue-Haul Greedy KV Migrate Only Replay Context Only Isolated Fastest

