


GPT-OSS-20B / A100; assumed shared WAN 400 Gbit/s; fixed-plan p05-p95 sensitivity
Independent pooled execution; fleet transfer assumptions apply

Shed power proxy (MW)



Queue-Haul LP Queue-Haul Greedy KV Migrate Only Replay Context Only Isolated Fastest