





QH LP KV only Isolated Fastest
 QH Greedy Replay only



Require predicted local queues to clear by deadline.
 2 MW installed/site; 1000 Gb/s shared WAN; eight GPUs share each measured VM endpoint.
 Source power: linear active-to-awake-idle allocation.
 One incoming pack/GPU; LP allocation except QH Greedy; queue clearance does not certify SLOs.