


3,920 input + 1,024 forced output tokens, 32 requests/point



—●— GPT-OSS-20B    -■- Qwen3.8-27B    -▲- Gemma-4-26B-A4B